Big data science Day 3

F. Legger - INFN Torino

https://github.com/Course-bigDataAndML/MLCourse-2223

Yesterday

- Big data, analytics
- Distributed computing
- ML: Feature engineering

Today

- **Machine learning**
  - Architectures
  - Train model and evaluate

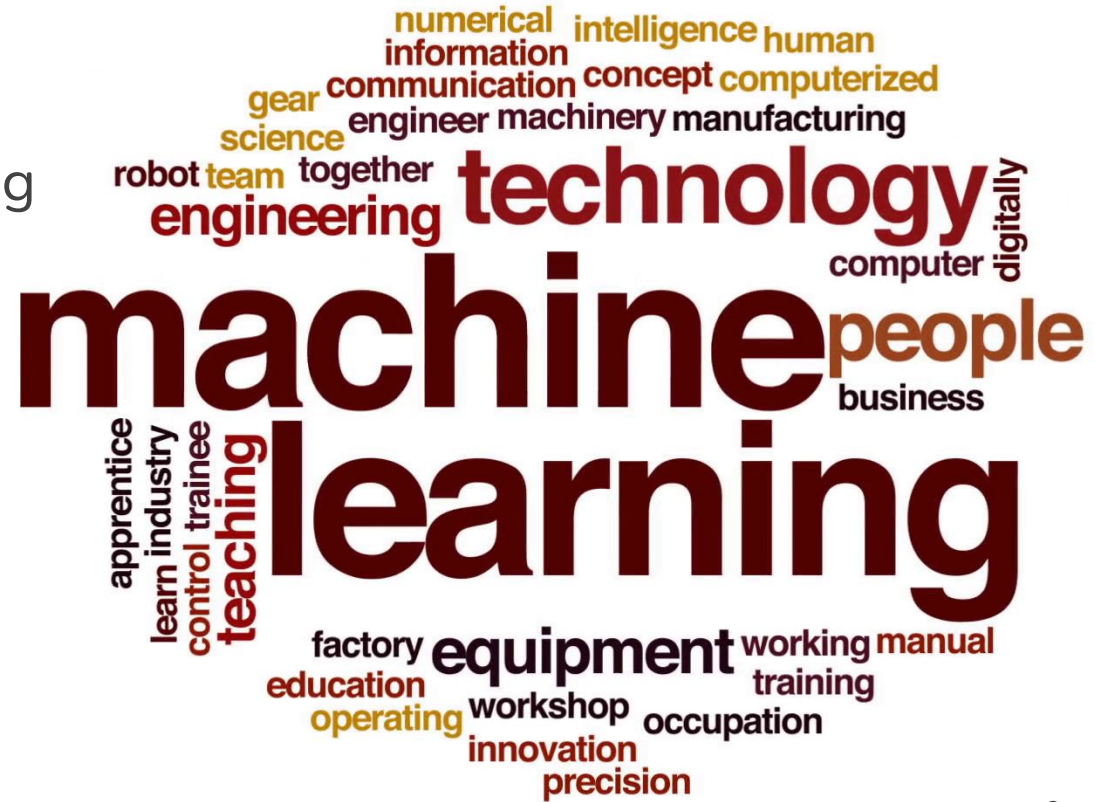

Remember:

Machine learning involves **two mathematical entities**

Model: a mathematical model describes the relationship between different aspects of the data

Features: a feature is a representation of raw data

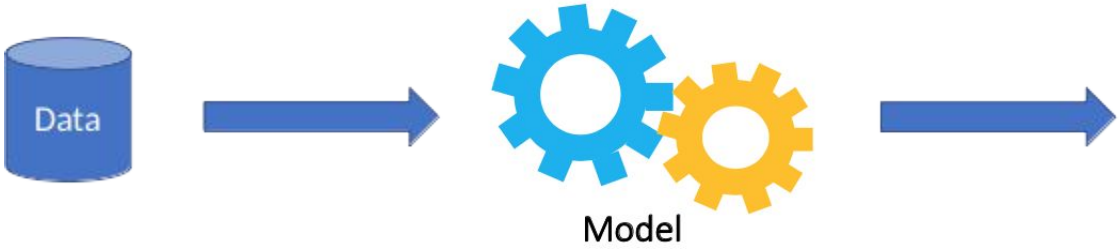

Feature engineering

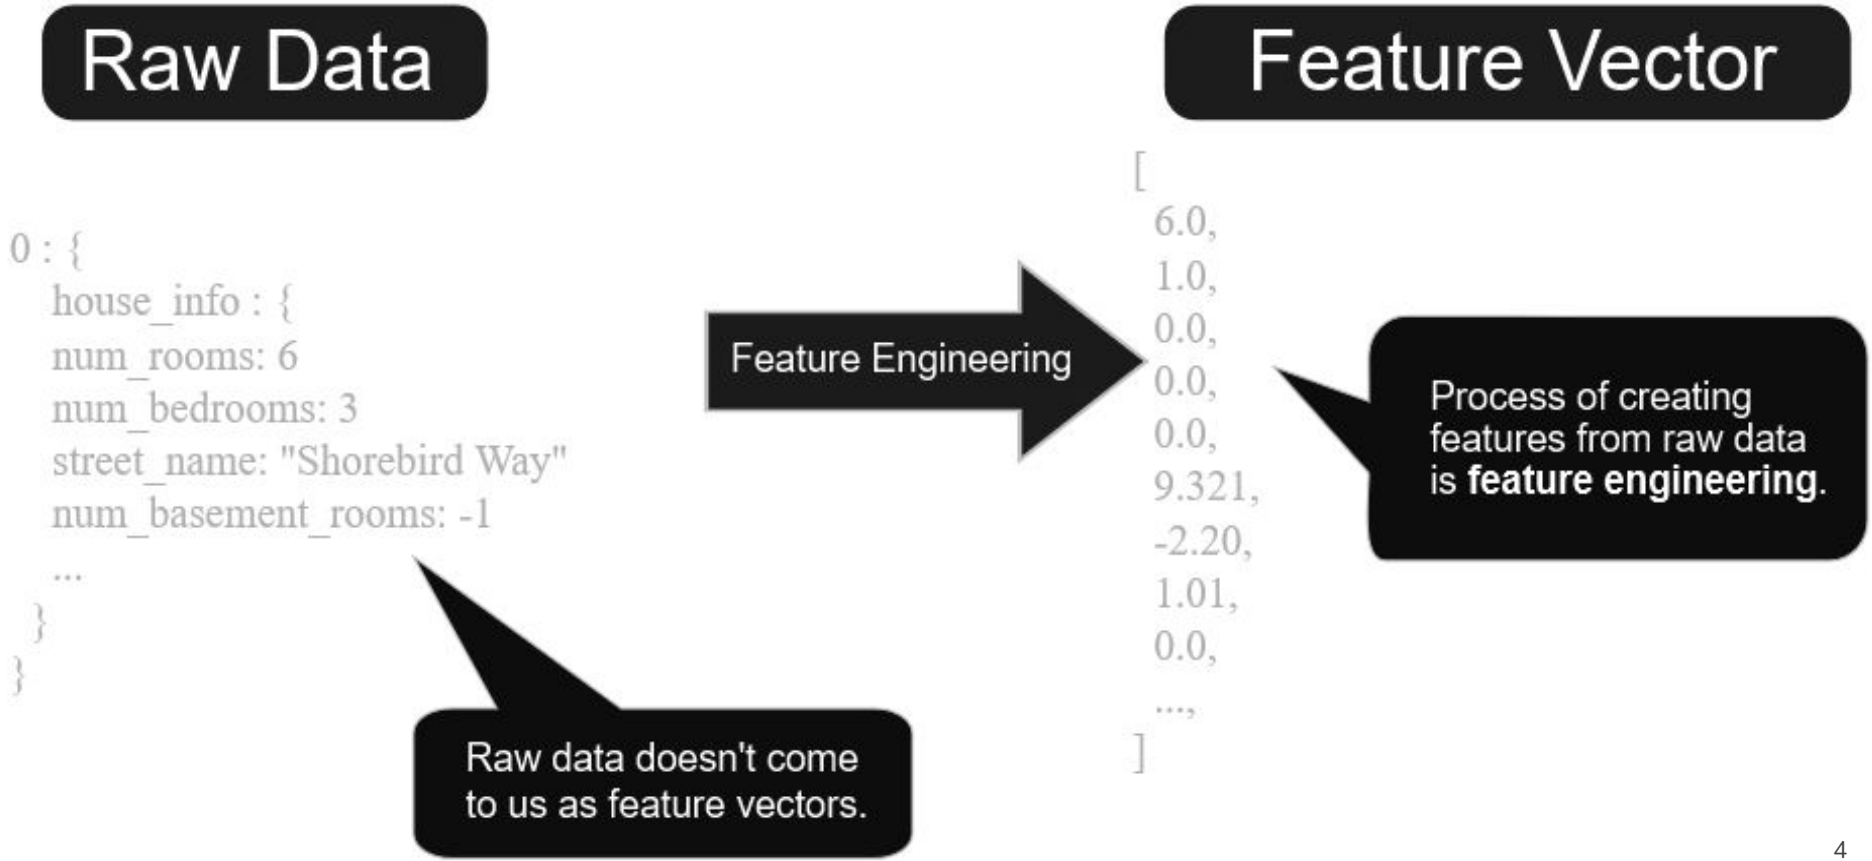

Example: supervised classification

Ingredients

- **Inputs: X**, is a matrix of size **n** (number of samples) x **m** (number of features)
- **Features: *X***, transformed inputs, matrix **nxm**
- **Labels: y,** vector size **n**

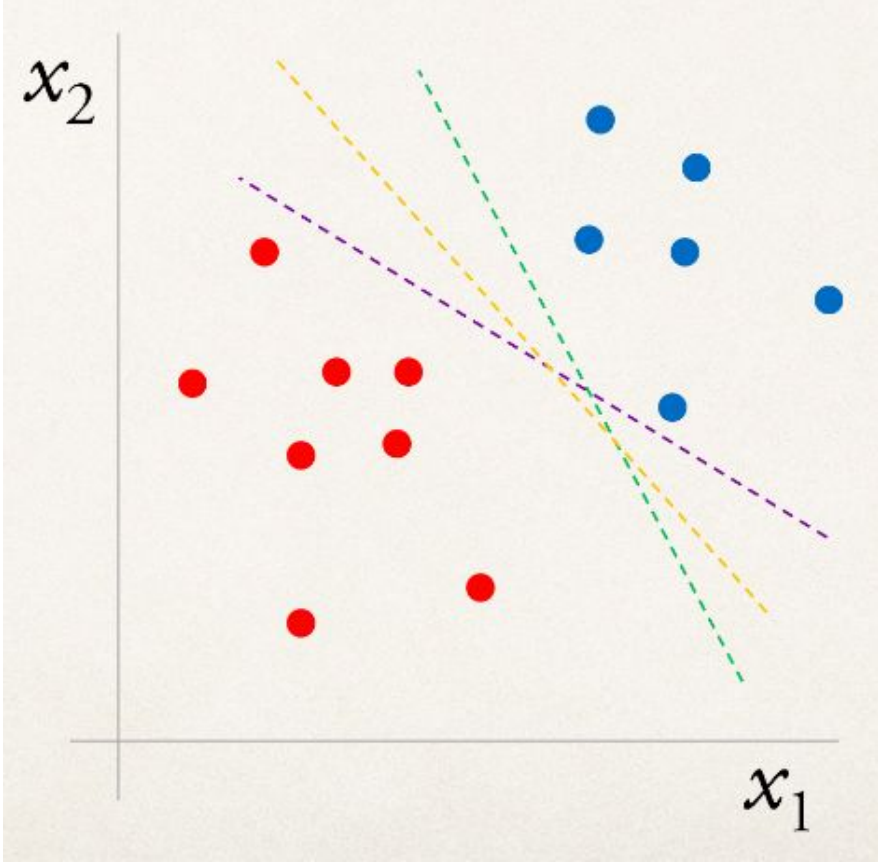

Recipe: supervised classification

- For each input vector X_i predict z_i, **i=1…n**

- **Predictions:**
  - $z_i = \varphi(W^T X_i)$ yields (0,1)
  - **Weights: W** (matrix) contains the model parameters
  - **Activation function: φ** (step function, sigmoid)

- **Cost function == loss function == prediction error**, function of the model parameters **W**
- Aim: find weights W that minimize cost function

Activation function

- Turns unbounded output into a known range/shape
- For example, **sigmoid** function only outputs numbers in the range (0, 1)
  - big negative numbers become ~0
  - big positive numbers become ~1.

$$S(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}.$$

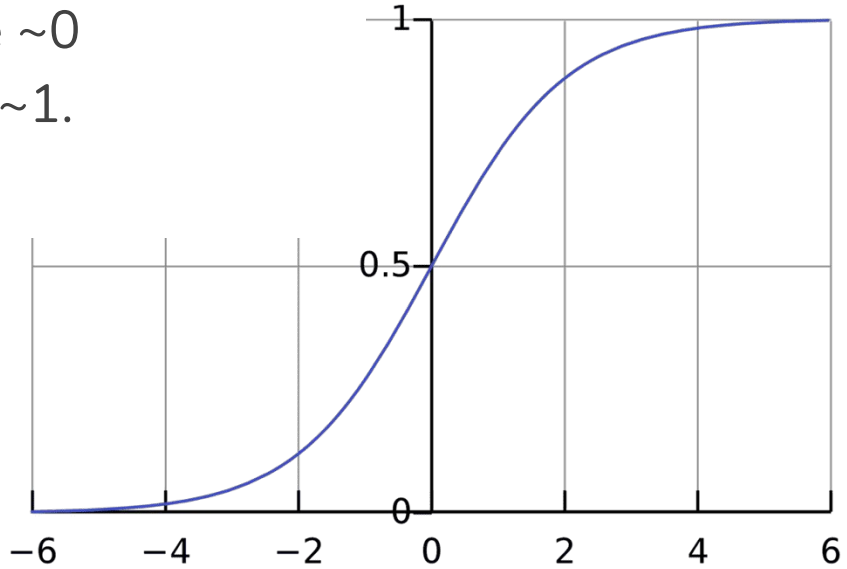

Another example, linear regression

- Inputs (features): $\mathbf{x_i}$
- Labels: $\mathbf{y_i}$

- Model: **y = a + bx**
- Weight+bias (parameters to be found): **a, b**
- Cost function: **Mean Square Error (MSE)**
- No **activation function:** problem is linear

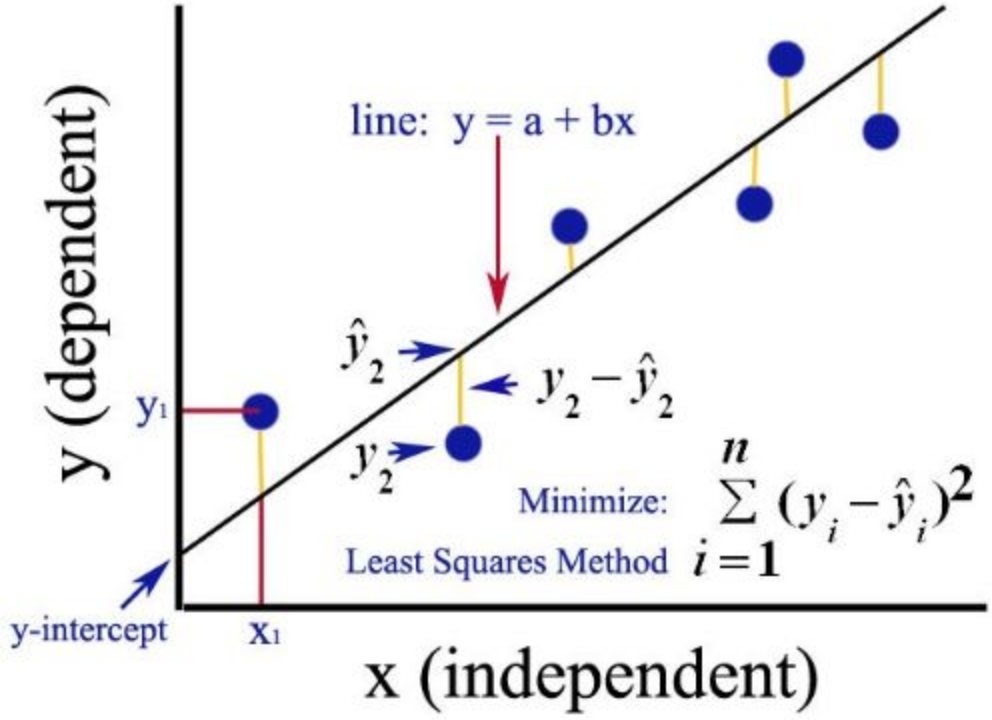

$$\mathrm{MSE} = \frac{1}{n}\sum_{i=1}^{n}(Y_i - \hat{Y}_i)^2.$$

Supervised learning

Decision trees: supervised classification

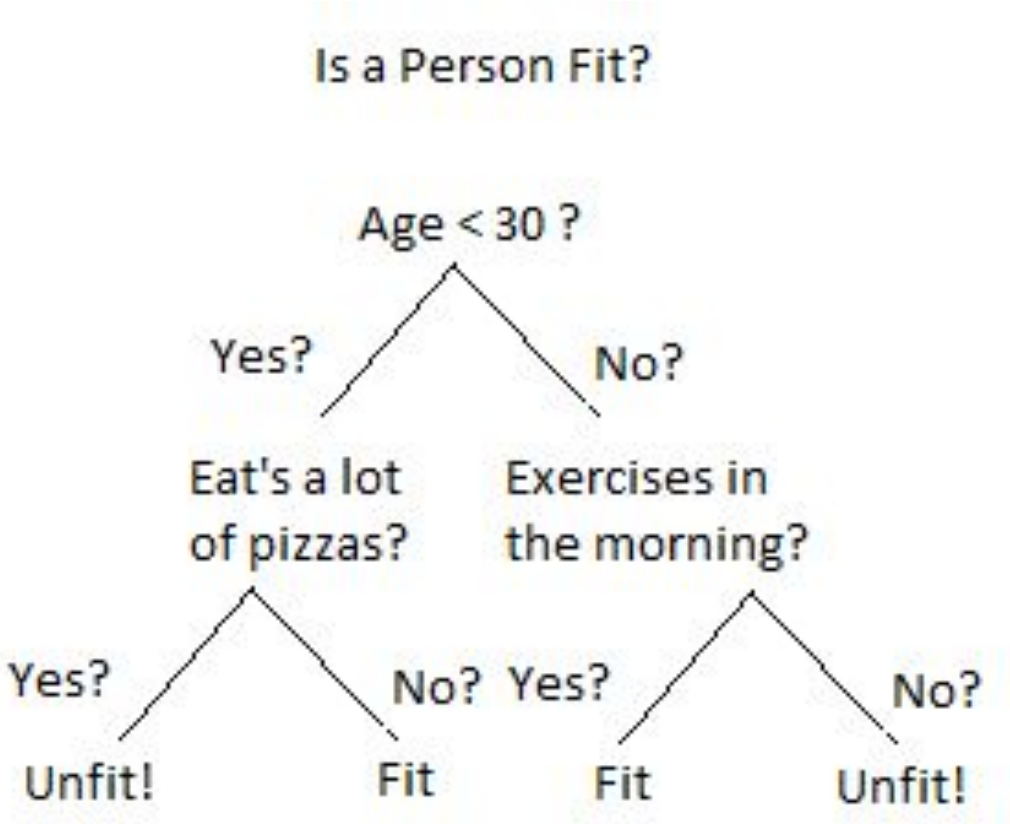

Typically used in combinations (Random forest, Gradient Tree Boosting)

Random Forest

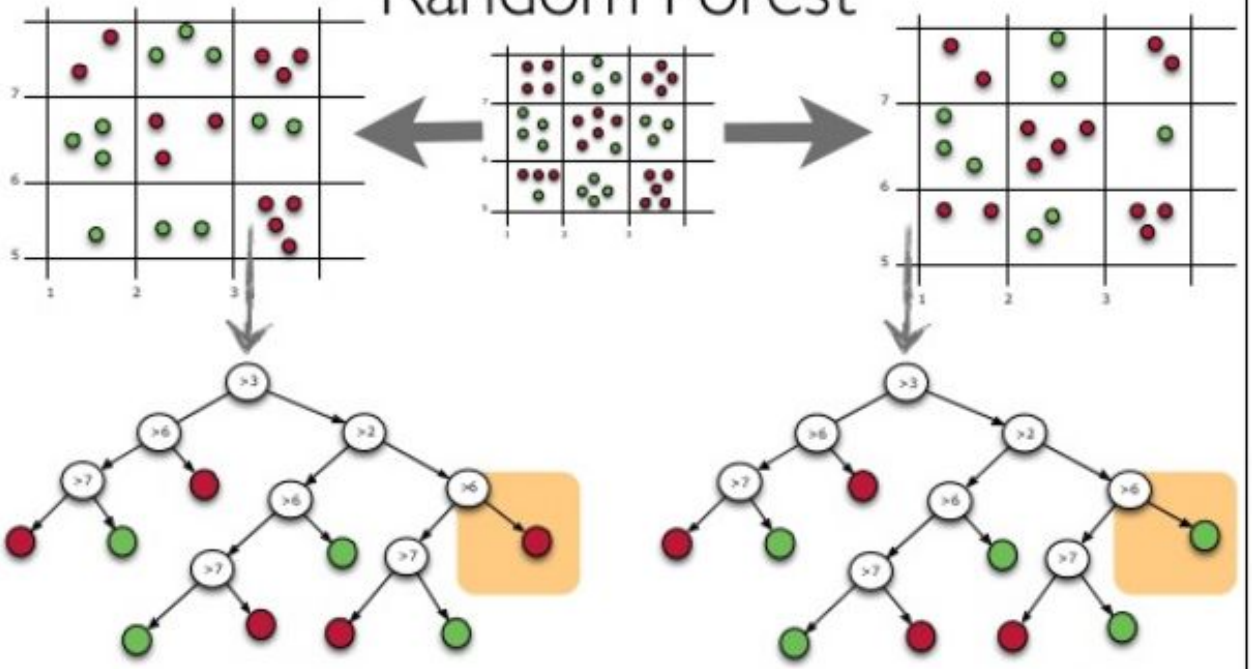

- Each tree sees part of the training sets and captures part of the information it contains

Ensembles

- **Bagging**
  - building multiple models (typically of the *same* type) from different subsamples of the training dataset
- **Boosting**
  - building multiple models (typically of the *same* type) each of which learns to fix the predictions errors of a prior model in the chain
- **Stacking**
  - building multiple models (typically of *different* types) and a supervisor model that learns how to best combine the predictions of the primary model
- **Weighting|Blending**
  - combine multiple models into single prediction using different weight functions

Neural networks: supervised classification

- Basic unit: **Neuron.** A neuron takes inputs, does some math with them, and produces **one** output

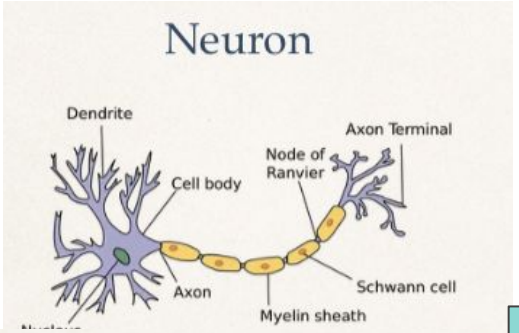

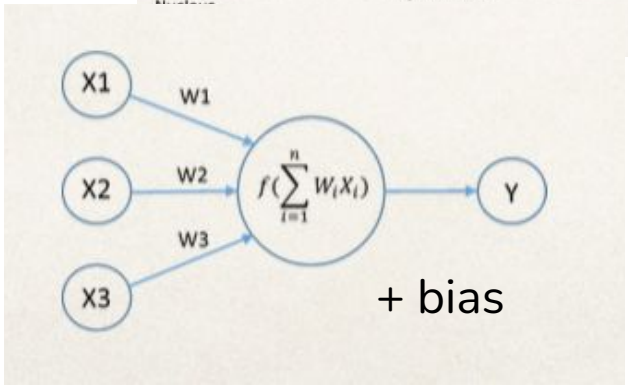

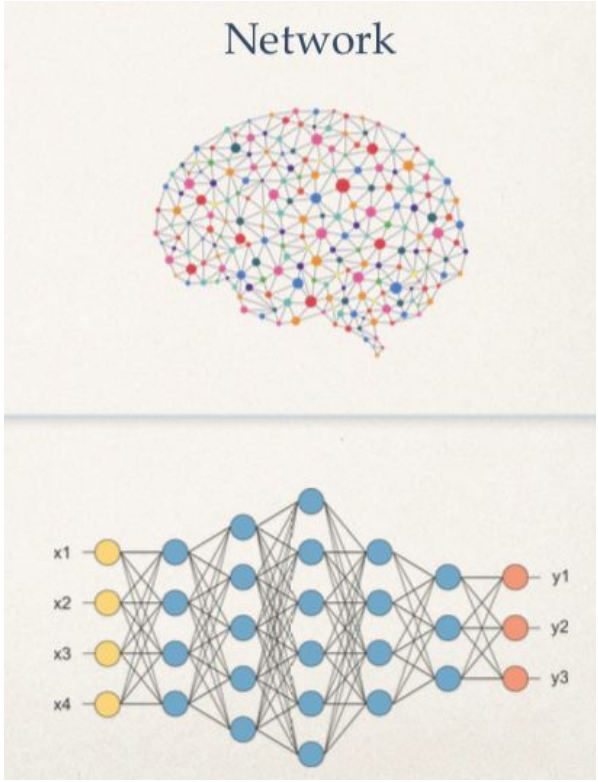

Support vector machines (SVG), supervised classification

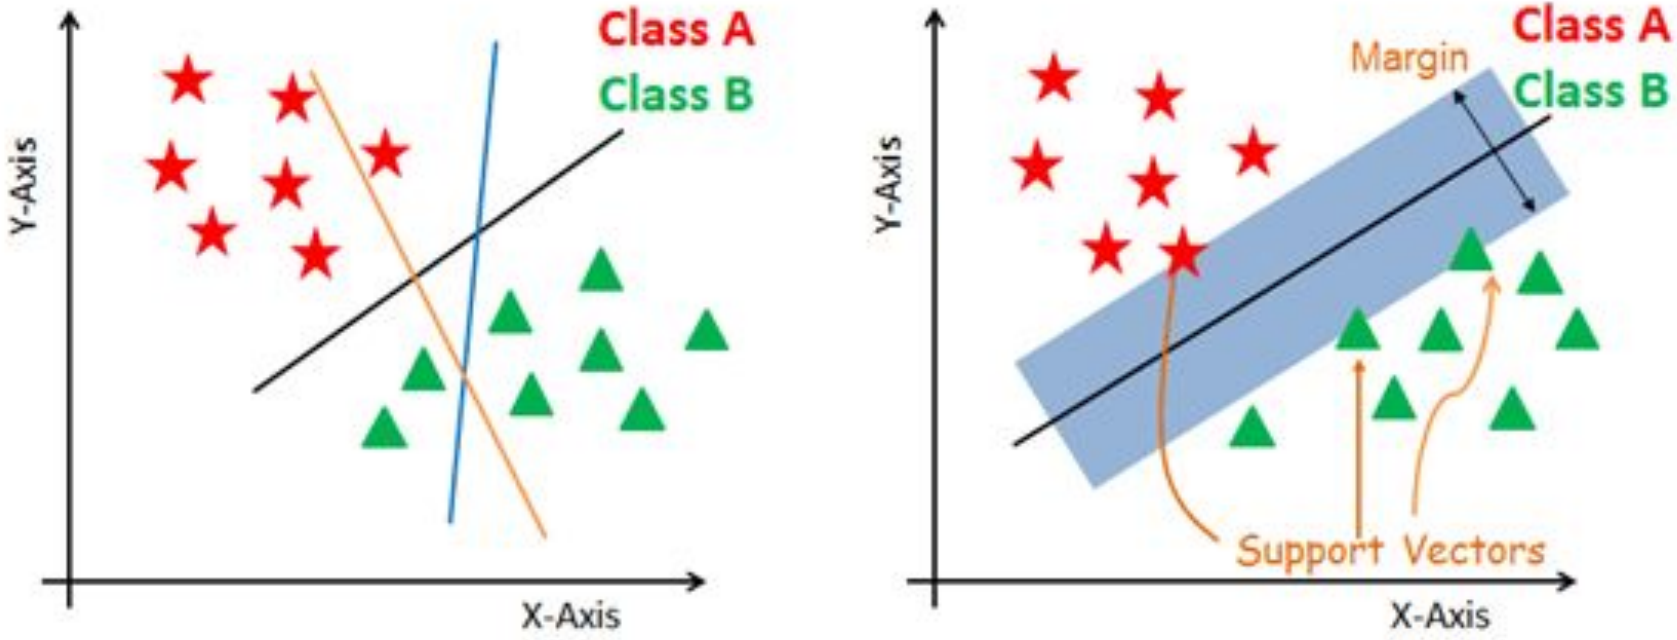

Unsupervised learning

Challenges

- No label (ground truth) in input dataset
- The system must have the ability to recognize patterns in the data without explicitly being told what patterns to identify

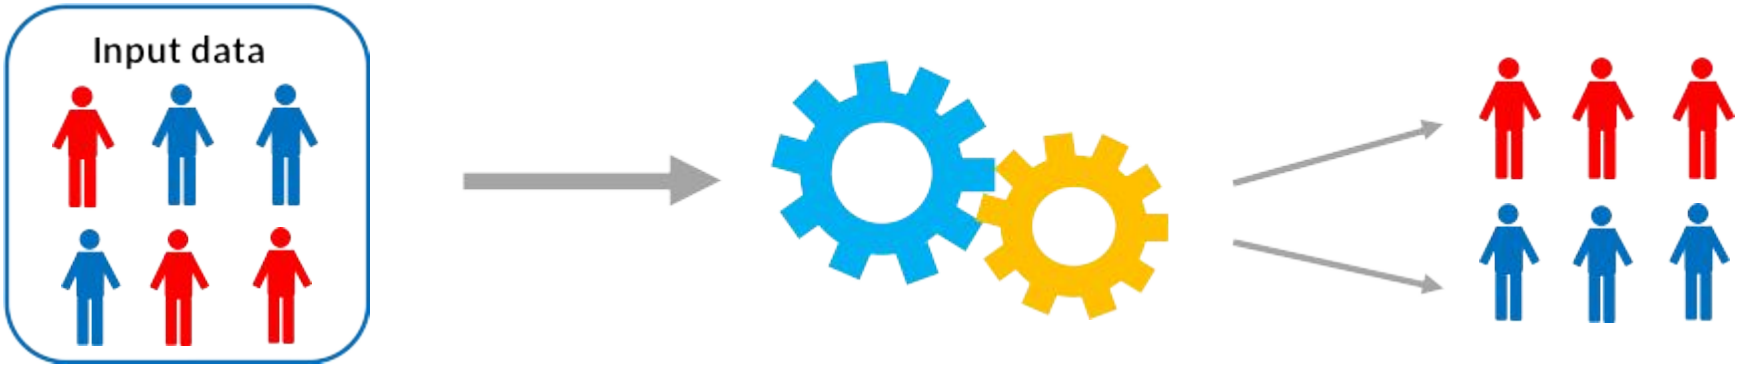

Principal Component Analysis (PCA), unsupervised

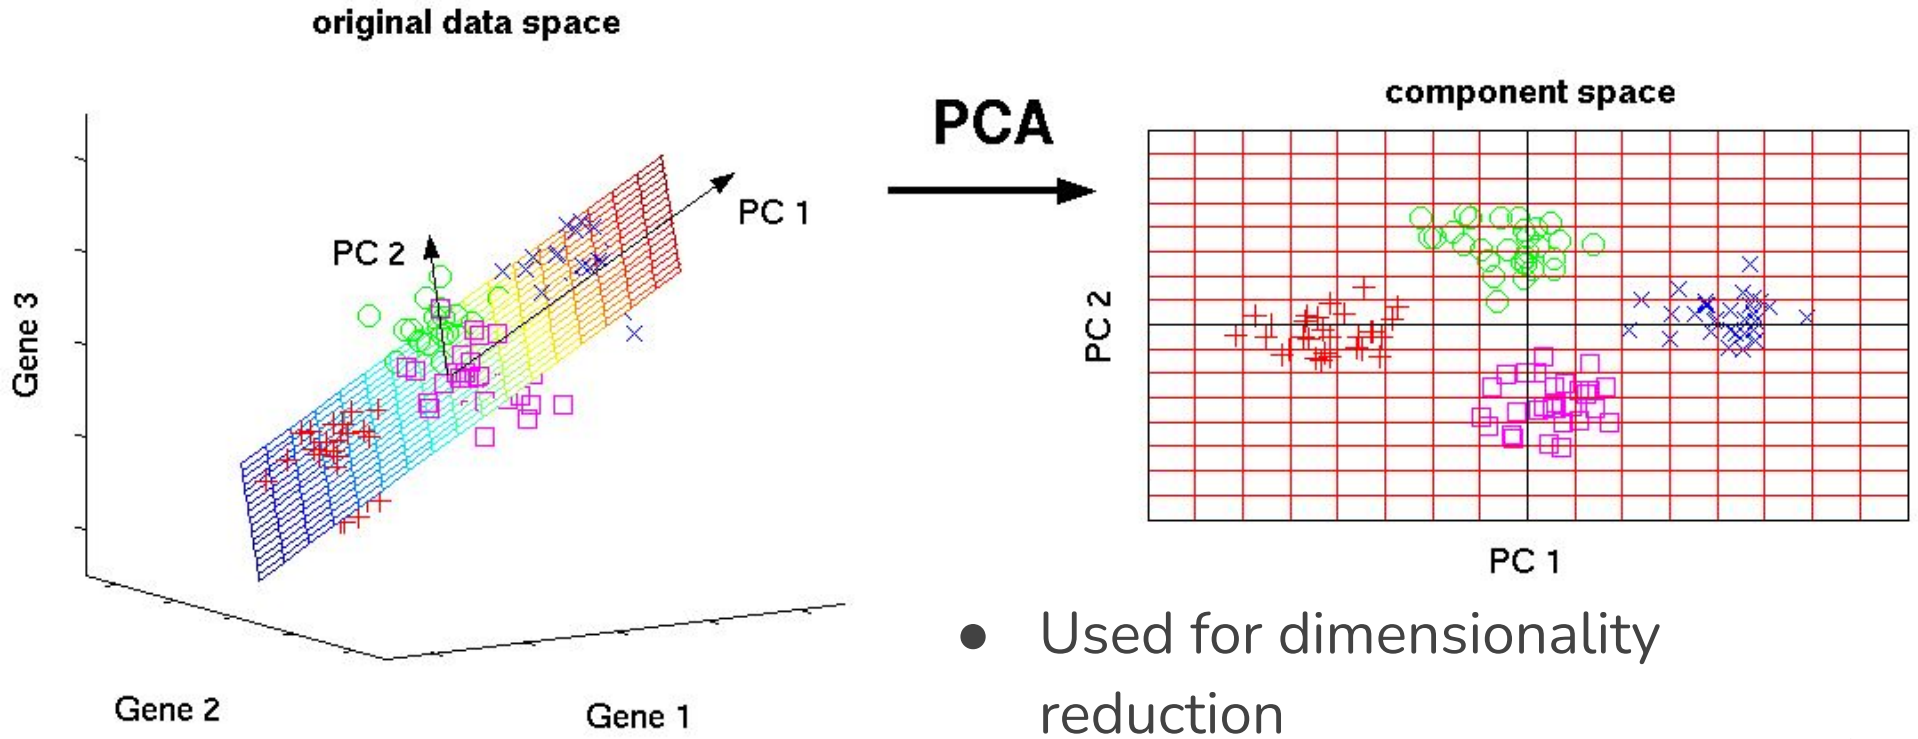

- Used for dimensionality reduction

K-means clustering, unsupervised

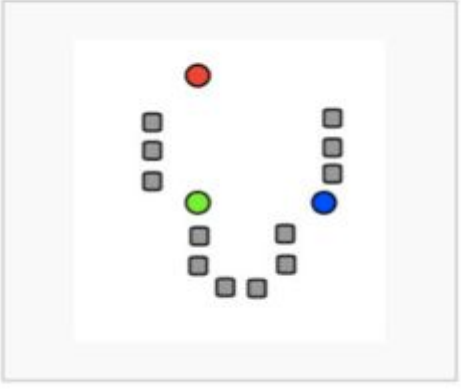

1. *k* initial "means" (in this case *k*=3) are randomly generated within the data domain (shown in color).

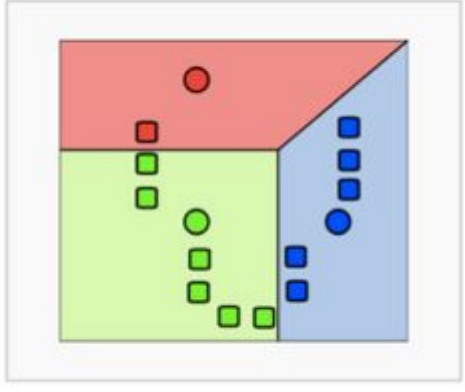

2. *k* clusters are created by associating every observation with the nearest mean. The partitions here represent the Voronoi diagram generated by the means.

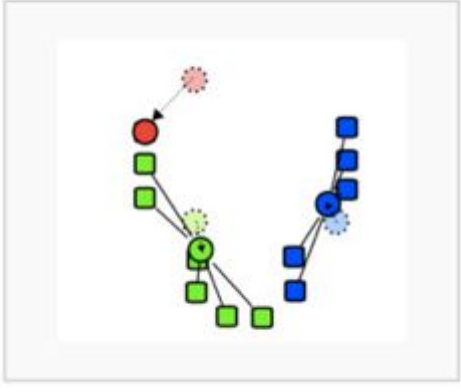

3. The centroid of each of the *k* clusters becomes the new mean.

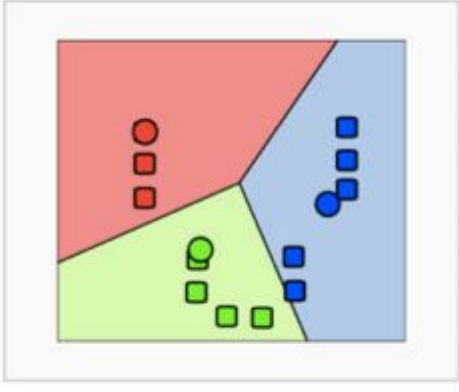

4. Steps 2 and 3 are repeated until convergence has been reached.

You must define **k, the number of clusters,** and which **distance** to use!

K-means: Distances

Taxicab or **Manhattan** distance: sum of the projections along all axis

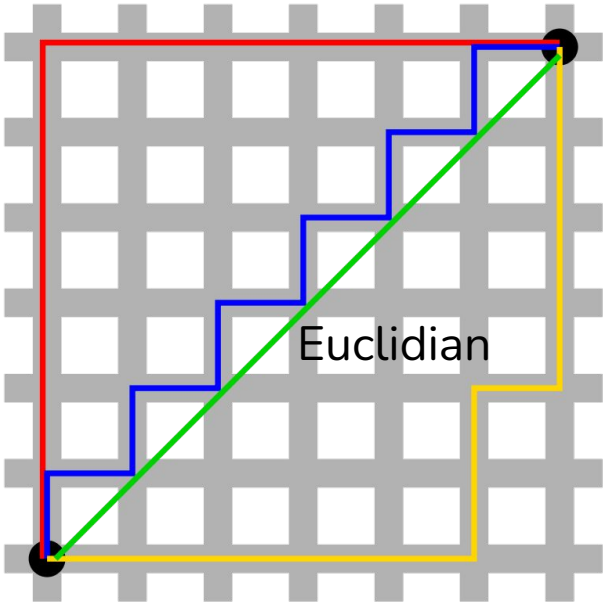

Euclidian distance

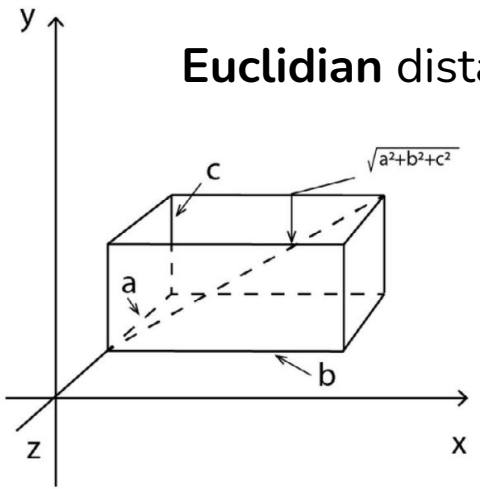

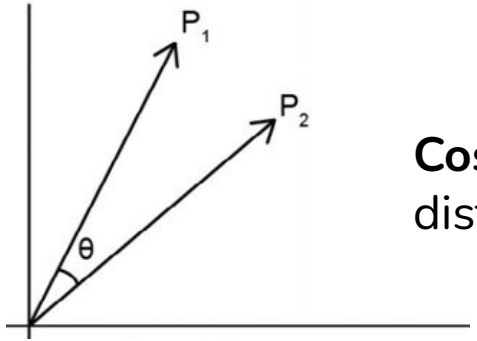

Cosine distance

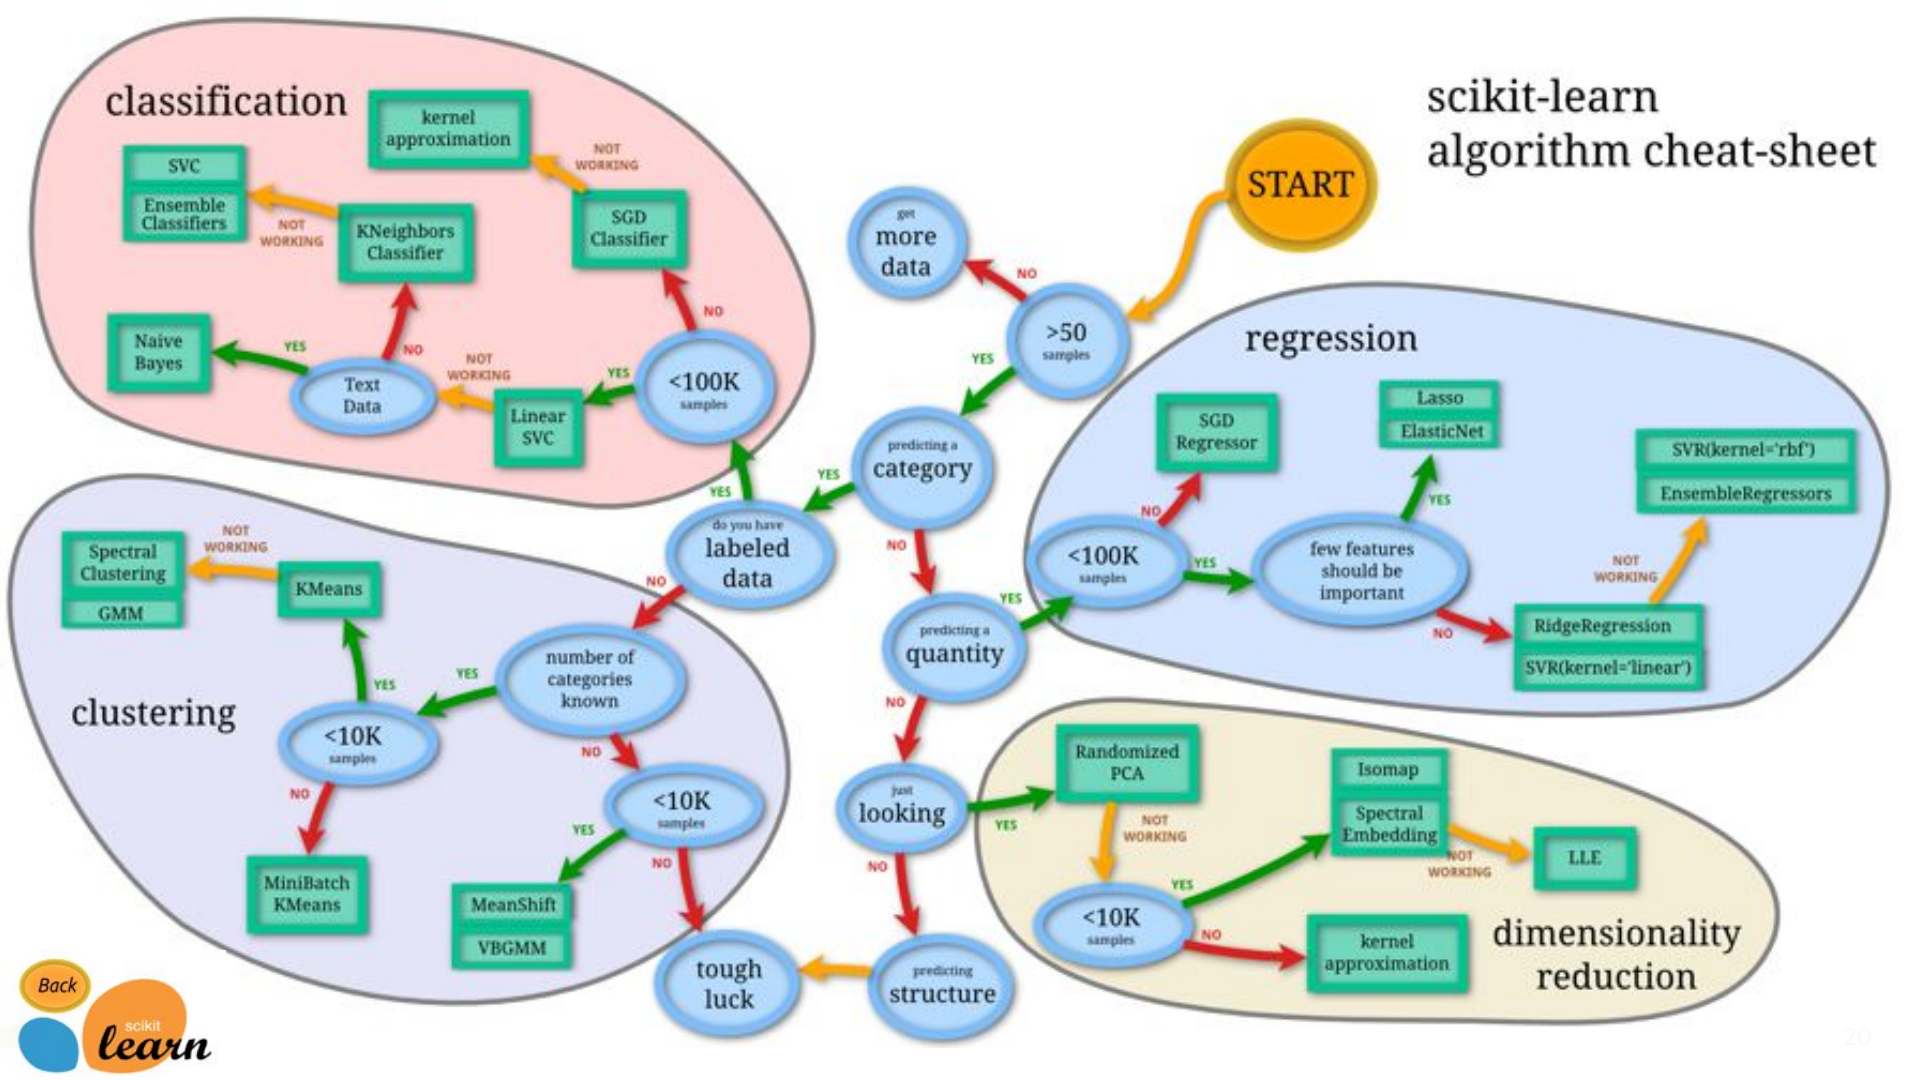
scikit-learn
algorithm cheat-sheet
START
get more data
NO
>50 samples
YES
classification
kernel approximation
NOT WORKING
SVC
Ensemble Classifiers
NOT WORKING
KNeighbors Classifier
SGD Classifier
NO
Naive Bayes
YES
Text Data
NO
NOT WORKING
Linear SVC
YES
<100K samples
YES
predicting a category
YES
NO
regression
SGD Regressor
Lasso
ElasticNet
SVR(kernel='rbf')
EnsembleRegressors
NO
<100K samples
YES
few features should be important
YES
NOT WORKING
NO
RidgeRegression
SVR(kernel='linear')
do you have labeled data
NO
predicting a quantity
YES
NO
clustering
Spectral Clustering
GMM
NOT WORKING
KMeans
YES
number of categories known
YES
NO
<10K samples
NO
MiniBatch KMeans
<10K samples
YES
NO
MeanShift
VBGMM
just looking
YES
NO
Randomized PCA
NOT WORKING
Isomap
Spectral Embedding
NOT WORKING
LLE
YES
<10K samples
NO
kernel approximation
dimensionality reduction
tough luck
predicting structure
Back
scikit learn

TYPE	NAME	DESCRIPTION	ADVANTAGES	DISADVANTAGES
Linear	 Linear regression	**The "best fit" line** through all data points. Predictions are numerical.	Easy to understand – you clearly see what the biggest drivers of the model are.	X Sometimes too simple to capture complex relationships between variables. X Tendency for the model to "overfit".
	 Logistic regression	The adaptation of **linear regression** to problems of classification (e.g., yes/no questions, groups, etc.)	Also easy to understand.	X Sometimes too simple to capture complex relationships between variables. X Tendency for the model to "overfit".
Tree-based	 Decision tree	A graph that uses a **branching method** to match all possible outcomes of a decision.	Easy to understand and implement.	X Not often used on its own for prediction because it's also often too simple and not powerful enough for complex data.
	 Random Forest	Takes the average of many decision trees, each of which is made with a sample of the data. Each tree is weaker than a full decision tree, but **by combining them we get better overall performance.**	A sort of "wisdom of the crowd". Tends to result in very high quality models. Fast to train.	X Can be slow to output predictions relative to other algorithms. X Not easy to understand predictions.
	 Gradient Boosting	Uses even weaker decision trees, that are increasingly **focused on "hard" examples.**	High-performing.	X A small change in the feature set or training set can create radical changes in the model. X Not easy to understand predictions.
Neural networks	 Neural networks	Mimics the behavior of the brain. Neural networks are interconnected neurons that pass messages to each other. Deep learning uses several **layers of neural networks put one after the other.**	Can handle extremely complex tasks - no other algorithm comes close in image recognition.	X Very, very slow to train, because they have so many layers. Require a lot of power. X Almost impossible to understand predictions.

All models are wrong, but some are useful (George Box)

Hyperparameters vs parameters

- Model **parameters** are learned during **training** when we optimize a loss function -> *weights*
- **Hyperparameters** are not model parameters and they **cannot** be directly trained from the data -> *model architecture*

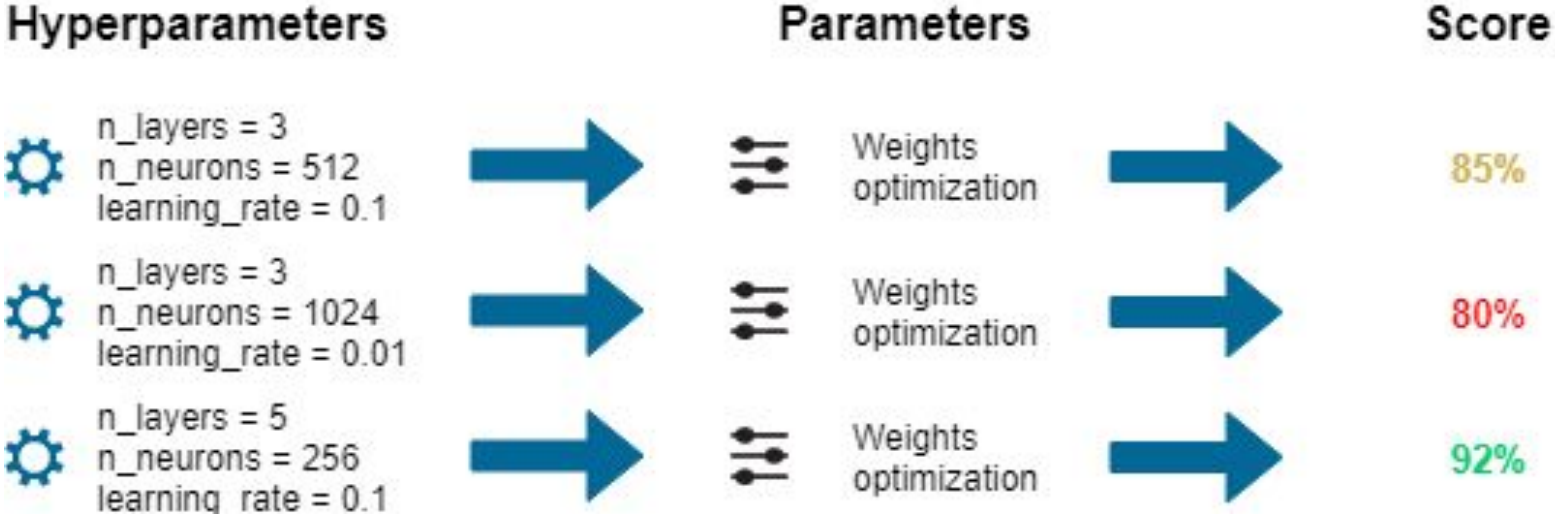

Hyperparameter tuning

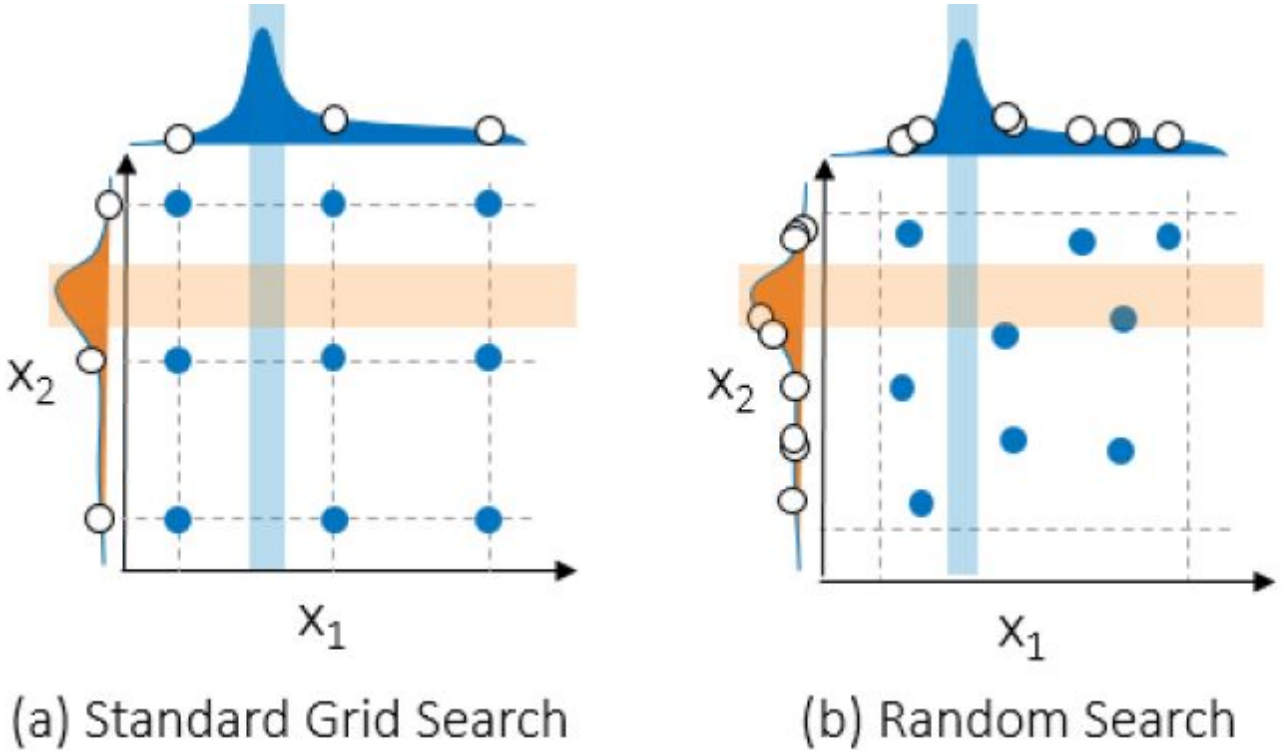

(a) Standard Grid Search

(b) Random Search

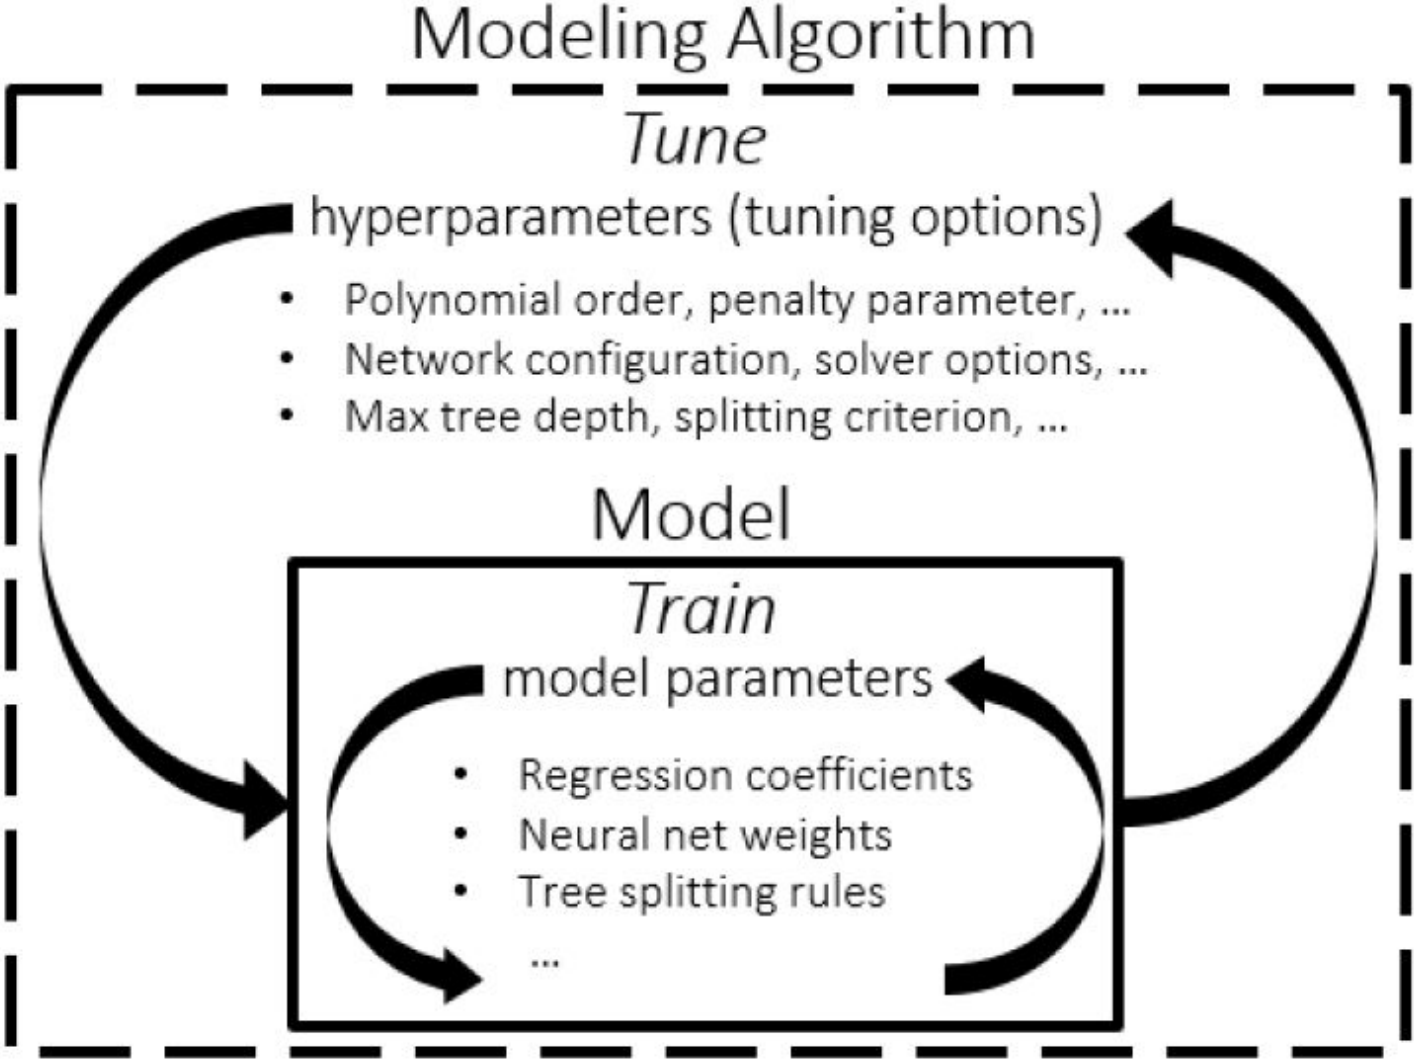
Modeling Algorithm
Tune
hyperparameters (tuning options)
• Polynomial order, penalty parameter, …
• Network configuration, solver options, …
• Max tree depth, splitting criterion, …
Model
Train
model parameters
• Regression coefficients
• Neural net weights
• Tree splitting rules
…

Overfitting / underfitting

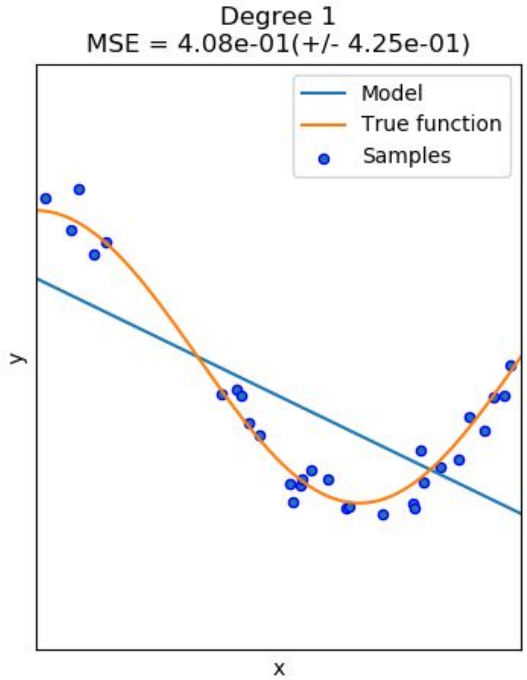

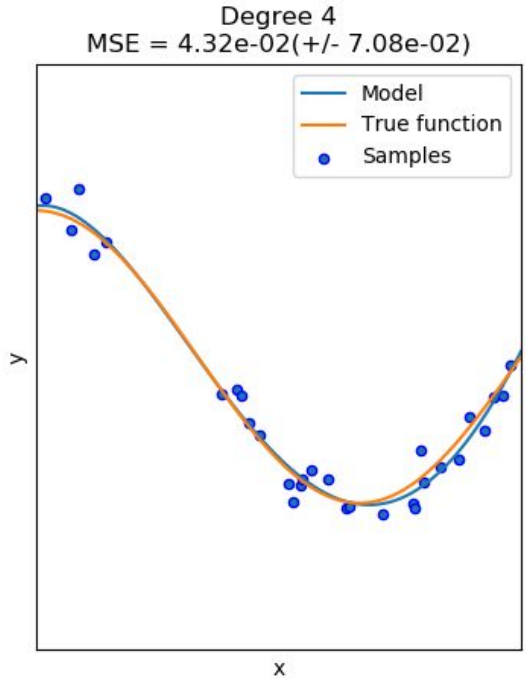

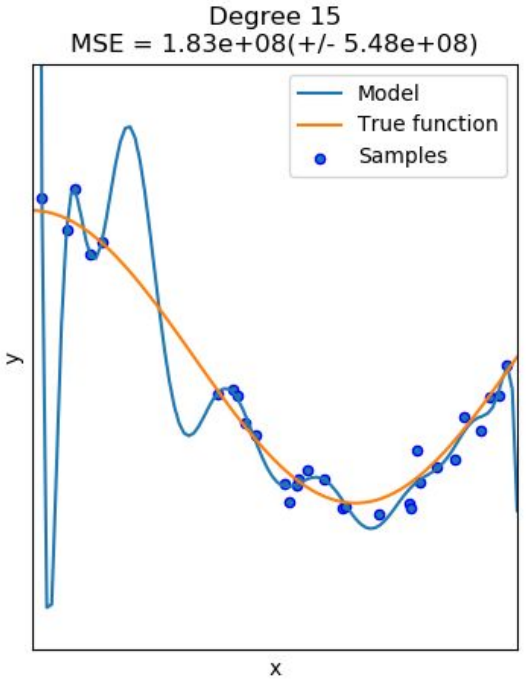

Underfitting
Model doesn't have enough (hyper-)parameters to describe data

Overfitting
Model has too many (hyper-)parameters

Classification metrics

- **ROC**: Receiver Operating Characteristics
- **AUC**: Area under the curve
- **TPR**: True positive rate
- **FPR**: False positive rate
- **TNR/FNR**: True/False negative rate

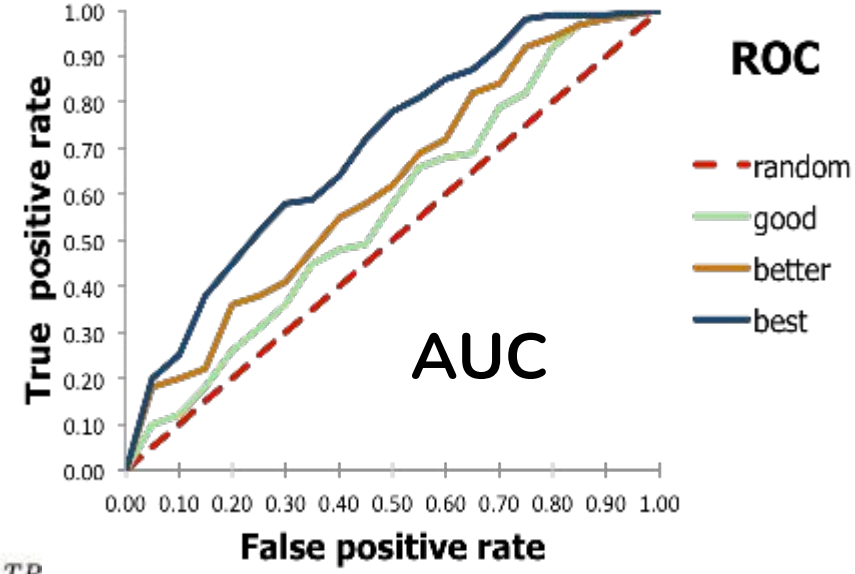

Hypothesized class	True class: p	True class: n
Y	True Positives	False Positives
N	False Negatives	True Negatives
Column totals:	P	N

$$\text{fp rate} = \frac{FP}{N} \qquad \text{tp rate} = \frac{TP}{P}$$

$$\text{precision} = \frac{TP}{TP+FP} \qquad \text{recall} = \frac{TP}{P}$$

$$\text{accuracy} = \frac{TP+TN}{P+N}$$

$$\text{F-measure} = \frac{2}{1/\text{precision}+1/\text{recall}}$$

Confusion matrix

Training and test set

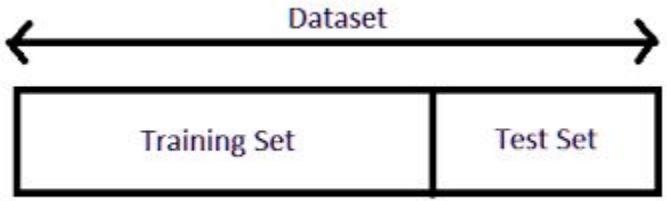

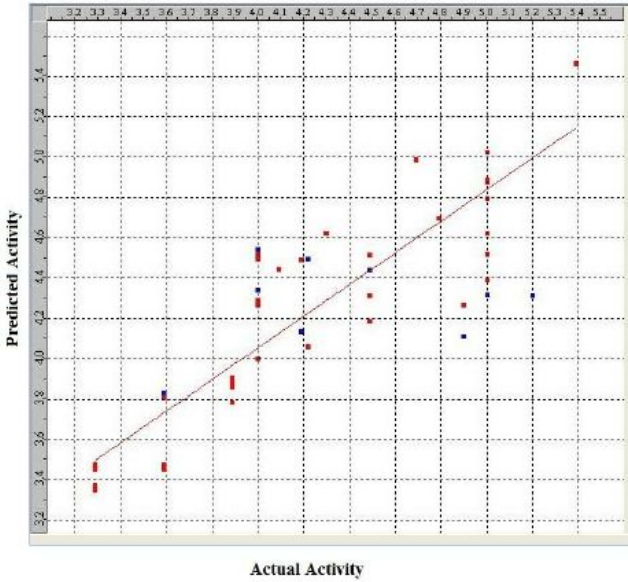

Cross-validation: is you model robust?

- Train/test split
- K-folds cross validation

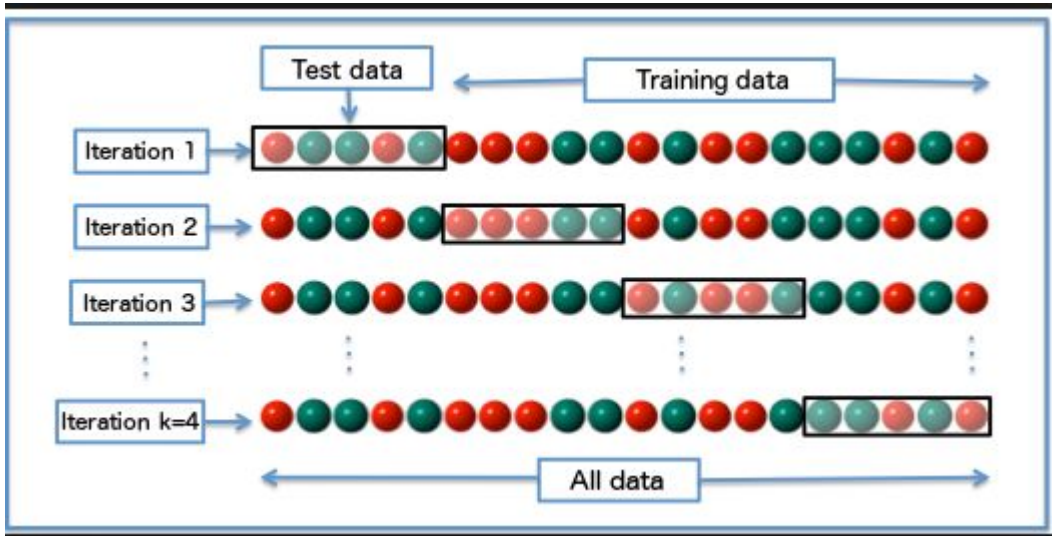

Loss, or could be any other metrics of interest

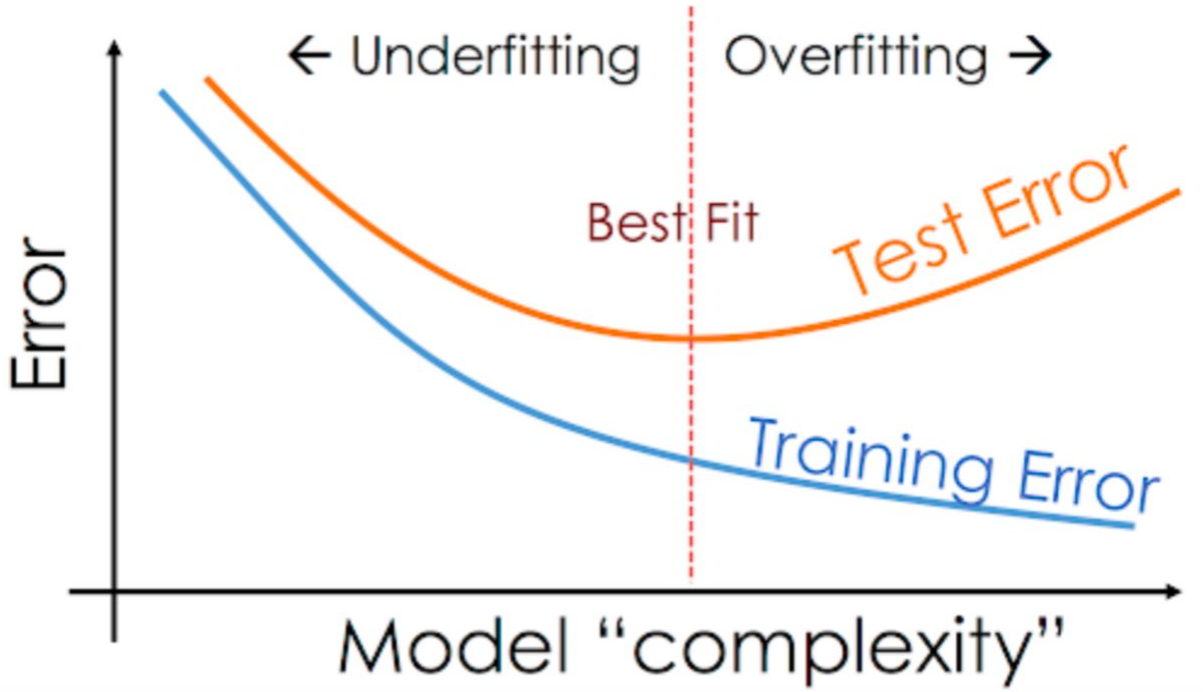

Or training epochs (number of iterations)

Hands-on today

1. **Point your browser to:** https://yoga.to.infn.it
2. **Open a terminal:**
    - cd MLCourse-2223
    - git pull
    - cp Notebooks/Day2/* ../
3. **From JupyterHub Home tab:**
    - start and run *ML_GBT.ipynb*

- Apache ML Library MLLib
    - Gradient Boosting Trees (GBT)
    - Hyperparameter optimisation
    - Multilayer Perceptron Classifier (MPC) - Bonus